

Daily Log

Monday February 10

I worked on creating the now called juggling environment because instead of sequentially training games, the agent has to play on each environment every 3 times. Everything but the function that runs an episode was successfully changed to handle a list of environments.

Wednesday January 12

I presented my syslab project and listened to other students' presentations

Friday February 14

I worked on saving and loading models to show the agent performing in the environment (rendering on). I debugged why the juggling environment wasn't working as expected. Can probably fix it in one class.

Timeline

Date	Goal	Met
Today minus 2 weeks	Get DIAYN Running	Not yet
Today minus 1 week	Get DIAYN Running	yes
Today	Benchmark DIAYN, discrete SAC, and DDQN on juggling	not yet
Today plus 1 week	Benchmark DIAYN, discrete SAC, and DDQN on triathlon	
Today plus 2 weeks	Make DBH (my algorithms name)	

Reflection

I can probably get juggling to work in a class period. The rest of the code to benchmark DIAYN, discrete SAC, and DDQN is ready (creating graphs and showing it working). I think I'll record the screen using obs to make gifs of the agent playing for presentations. I wonder if also letting people play against the agent would be interesting. Each agent is its own difficulty and maybe there's a point where people can't seem to beat it at pong. Not sure how pong has 6 discrete actions possible, but I can probably make some converter.

I think I have around a month to put my algorithm, but it should be enough. At the very least I will have something to show.